

Continuous_Control

May 10, 2022

0.1 The learning algorithm

This reinforcement learning task was solved using Deep Deterministic Policy Gradient, an algorithm created by OpenAI. DDPG is closely related to Q-learning but adapted for use in continuous spaces. The algorithm utilises innovations from Q-learning such as fixed target and experience replay. The algorithm falls under an umbrella called the Actor-Critic method and consists of 4 networks: actor, critic, target-actor and target-critic.

The algorithm uses the current actor-network to choose an action a based on input state s , and the environment returns reward r and next_state s' . In the replay buffer, this is stored as an experience tuple, (s, a, r, s') . The replay buffer is used to break correlations between consecutive experience tuples. The max size is $1e6$, and once it is full, the oldest experiences are discarded. Once `BATCH_SIZE` experiences are in the buffer, the agent begins learning.

The current actor and critic networks are updated as follows: 1) A minibatch of experience tuples is sampled from the replay buffer randomly. 2) Action a' is then chosen by the target-actor and evaluated by the target-critic. 3) This value is then discounted by γ , added to reward r to form y . 4) The loss is then calculated by the MSE of y , and the expected value of action a calculated by the current critic. 5) This loss is used to update the current critic through backpropagation. 6) We then update the actor-network by taking the derivative of the critic network with respect to the policy parameters, using the mean of the gradients in the mini-batch.

The target networks are just delayed copies of the actor and critic networks, which improves stability as the agent is chasing a much slower-moving target. Every episode, these target networks are updated by a small amount, τ , to match the current actor and critic closely. This project performs a hard update (current weights completely transferred to target weights) every `HARD_FREQ`, reducing training time without causing any stability issues.

0.2 Hyperparameter Choice

```
BUFFER_SIZE = int(1e6) BATCH_SIZE = 512 # chosen to fit in GPU memory
GAMMA = 0.90 # discount factor
TAU = 0.5e-3 # low due to addition of hard updates
LR_ACTOR = 1e-4 # learning rate of the actor
LR_CRITIC = 1e-4 # learning rate of the critic
WEIGHT_DECAY = 0 # L2 weight decay
EPSILON_DEC = 0.997 # noise reduction
TRAIN_FREQ = 10 # how often train agent
HARD_FREQ = 100 # how often to perform the hard update
TRAIN_N = 5 # how many times train agent
```

Network architecture is the same for all nn was inspired by Gkwalik.

```
fc_1 = 350
fc_2 = 280
```

0.3 Enviroment Setup

```
In [1]: !pip -q install ./python
        from unityagents import UnityEnvironment
        import numpy as np

        env = UnityEnvironment(file_name='/data/Reacher_One_Linux_NoVis/Reacher_One_Linux_NoVis.

        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatib
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 3.0.
jupyter-console 6.4.3 has requirement jupyter-client>=7.0.0, but you'll have jupyter-client 5.2.
```

```
INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
    Number of Brains: 1
    Number of External Brains : 1
    Lesson number : 0
    Reset Parameters :
        goal_speed -> 1.0
        goal_size -> 5.0
Unity brain name: ReacherBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 33
    Number of stacked Vector Observation: 1
    Vector Action space type: continuous
    Vector Action space size (per agent): 4
    Vector Action descriptions: , , ,
```

0.4 Imports

```
In [2]: import gym
        import random
        import torch
        import numpy as np
        from collections import deque
        import matplotlib.pyplot as plt
        %matplotlib inline
```

```

from Agent import Agent

env_info = env.reset(train_mode=True)[brain_name]

agent = Agent(state_size=33, action_size= 4, random_seed=2)

```

0.5 Training Loop

```

In [3]: def ddpq(n_episodes=1000, max_t=1000, print_every=100):
    scores_deque = deque(maxlen=print_every)
    scores = []
    for i_episode in range(1, n_episodes+1):
        env_info = env.reset(train_mode=True)[brain_name]
        state = env_info.vector_observations[0]
        agent.reset()
        score = 0
        for t in range(max_t):
            action = agent.act(state)
            env_info = env.step(action)[brain_name]
            next_state = env_info.vector_observations[0]
            reward = env_info.rewards[0]
            done = env_info.local_done[0]
            agent.step(state, action, reward, next_state, done)
            state = next_state
            score += reward
            if done:
                break
        scores_deque.append(score)
        scores.append(score)
        if i_episode % print_every == 0:
            print('\rEpisode {} \t Score: {:.2f} \t Average Score: {:.2f}'.format(i_episode,
                                        score, np.mean(scores_deque)))

            if np.mean(scores_deque) >= 30.0:
                print('\nEnvironment solved in {:d} episodes! \t Ended in episode: {:d}'.format(
                    i_episode, i_episode))
                torch.save(agent.actor_local.state_dict(), 'checkpoint_actor.pth')
                torch.save(agent.critic_local.state_dict(), 'checkpoint_critic.pth')
                break

    return scores
scores = ddpq()

```

```

/opt/conda/lib/python3.6/site-packages/torch/nn/functional.py:1795: UserWarning: nn.functional.tanh
  warnings.warn("nn.functional.tanh is deprecated. Use torch.tanh instead.")

```

Episode 100	Score: 7.97	Average Score: 2.88
Episode 200	Score: 36.53	Average Score: 17.67

Environment solved in 143 episodes!

Ended in episode: 243

0.6 Plot of Rewards

Solved after 243 episodes

```
In [9]: plt.plot(scores)
plt.xlabel('episode')
plt.ylabel('score')
plt.title('Plot of score over episodes')
plt.show()
```



0.7 Future Improvement

GAMMA, TAU and HARD_FREQ can be tuned further. In addition, it is likely that placing priority on important/rare experience tuple will lead to faster convergence. Also a lot more investigation can be done on the exploration vs exploitation front, as adding noise is far more abstract than the epsilon greedy policy used in regular Q-learning. In addition, the Ornstein-Uhlenbeck process may be an overly complex way to add noise, and simpler methods like normally distributed may also work.